



भारतीय प्रौद्योगिकी संस्थान जोधपुर
Indian Institute of Technology Jodhpur

Clustering and PageRank

Assignment-IV

Machine Learning with Big Data

Submitted by

Shreyas B Gaikwad

Roll No. M25DE1042

M Tech (Data Engineering)

GITHUB LINK : https://github.com/geek1042/mlbd_a4.git

INTRODUCTION

The assignment covers distributed clustering algorithms, inverted index construction for web search and graph-based ranking using the PageRank algorithm on Apache Spark.

The dataset for Part 1 is the UCI Spambase dataset containing 4601 emails represented as 58 dimensional feature vectors. For Part 2 a collection of webpage text files was provided along with a set of search queries and expected answers. For Part 3, two graph files were used and a small graph and a whole graph with 1000 nodes and 8192 edges.

All three parts were implemented and executed on Google Colab due to PySpark giving error on my legacy system Sony Vaio (2012). PySpark 4.0.2 was used as the distributed computing framework throughout. For Part 3, an initial implementation that ran the full PageRank iteration loop inside Spark RDDs resulted in runtimes exceeding 15 minutes even on the smaller graph, due to task scheduling overhead on a single-machine session with limited RAM. The dataset files were uploaded directly to the Colab session storage and accessed from there.

<u>Part</u>	<u>Dataset</u>	<u>Size</u>
1	UCI Spambase	4601 pts $\times$ 58 dims
2	Webpage collection	5 pages, 11 search actions
3	Graph Files, small.txt, whole.txt	1000 nodes, 8192 edges.

SOLUTIONS

Part 1: Clustering

The goal was to cluster a 4601-point, 58-dimensional spam detection dataset into k groups using two different center-selection strategies.

PySpark was initialized first since the assignment uses MLlib Vectors as the data representation.

The `readVectorsSeq` function reads the dataset line by line and converts each comma-separated row into a dense PySpark Vector. This gives us a standard list of 4601 points, each with 58 features, which is the input format both clustering algorithms expect.

The `kcenter` function implements the Farthest-First Traversal algorithm. It picks one random starting center, then repeatedly selects the point that is farthest from all already-chosen centers.

The `kmeansPP` function improves on random initialization by sampling each new center with probability proportional to its squared distance from the nearest existing center.

The `kmeansObj` function computes the average squared distance of every point to its nearest center, the standard k-Means objective. A lower value means the centers explain the data better.

Assumption : `kmeansObj` computes the total sum / $|P|$ as defined. The output from the notebook is:

- `kcenter(P, k=10)` ran in **0.2481 seconds**
- `kmeansPP(P, k=10)` gave `kmeansObj` = 133,742,683.59
- The **coreset** approach — `kcenter(P, k1=50)` followed by `kmeansPP(X, k=10)` — gave `kmeansObj` = 5,149,794,765.93

The coreset objective is much higher than direct k-Means++ because 50 farthest first points. Running k-Means++ on just those 50 boundary points produces centers that represent the edges of the data well but miss the dense inner regions. This shows that the coreset quality depends heavily on k_1 , a larger k_1 would bring more inner dense points into the coreset and improve the final objective.

Part 2: Web Search — Inverted Index

The inverted index flips the usual document to word direction and not storing what words a document contains, it stores for each word which documents contain it and at exactly which positions.

Assumptions made: Stop words and punctuation are exactly the lists given in the problem. Word positions are 1 index and stop words are counted in position numbering but not stored. The plural to singular mapping only covers the pairs listed for it.

The tokenizer replaces punctuation with spaces, splits the text, lowercases everything, maps plurals to singulars and assigns each token a position including stop words which given a position but are not in index. This ensures position numbers remain accurate even when stop words are skipped.

The seven classes build up the index layer by layer:

- `Position` stores a (page, index) pair.
- `WordEntry` collects all positions for one word across all documents and computes term frequency.
- `PageIndex` is a per-page word-to-positions map.
- `PageEntry` reads a file, tokenizes it, and builds its `PageIndex`
- `MyHashTable` implements rolling hash for collision resolution
- `InvertedPageIndex` ties everything together by mapping every word across all pages to its full position list.
- `SearchEngine` parses action strings and dispatches to the right query.

Output matched `answers.txt` exactly, all 11 lines of output were identical, correct handling of missing words, missing pages, position reporting, and plural normalization was performed.

Part 3: PageRank

PageRank assigns each node a score representing how likely a random surfer is to land there. The surfer follows outgoing links with probability $\beta = 0.8$, and teleports to any random node with probability 0.2. Starting from uniform scores, this is updated iteratively for 40 rounds.

Assumption: The dataset contains nodes with no outgoing edges (dangling nodes), despite the problem statement saying all nodes have positive out-degrees. The small graph loaded as 100 nodes with 950 edges rather than 53 nodes, which suggests the `small.txt` file used contained a different graph than described — but the algorithm was verified by checking that the top score fell near 0.036 and the rank sum equalled 1.0.

Spark was used to load and deduplicate edges via `.distinct()` on an RDD, satisfying the requirement to process `M` as an RDD. The actual iterative computation was done using NumPy matrix-vector multiplication, because running RDD joins inside a 40-iteration loop on a 1000-node graph was causing 15+ minute runtimes due to Spark task scheduling overhead far exceeding the actual computation cost.

Dangling nodes, those that only appear as destinations and have no outgoing edges caused rank to leak out each iteration, resulting in a rank sum of 0.784 instead of 1.0 and a wrong top score of 0.026. The fix redistributes the total rank held by dangling nodes uniformly to all nodes at each and every iteration, which is the standard Google Matrix correction and restores rank conservation.

Final outputs:

Small graph — top score = 0.035731 (expected ≈ 0.036), sum = 1.0

Whole graph: Top 5 nodes: 263, 537, 965, 243, 285

Bottom 5 nodes: 408, 424, 62, 93, 558

sum = 1.0

CODE AND OUTPUT (uploaded to github, https://github.com/geek1042/mlbd_a4.git)

✓ CSL7110

MACHINE LEARNING WITH BIG DATA

Assignment 4: Clustering and PageRank

Shreyas Gaikwad M25DE1042 M Tech (Data Engineering)

GitHub Repository: https://github.com/geek1042/mlbd_a4.git

Dataset Source: Uploaded to Colab directly.

Part 1: Clustering

Farthest-First Traversal (k-center): This is a greedy approximation algorithm for the k-center problem. Starting from a random point, it repeatedly picks the point that is farthest from the already-chosen centers. It guarantees a 2-approximation to the optimum solution and runs in $O(|P| \times k)$ time since for each of k iterations, so we scan all |P| points to find the maximum distance.

k-Means++: k-Means++ is a smarter initialization strategy for k-Means. Instead of choosing centers uniformly at random, it chooses each successive center with probability proportional to its squared distance from the nearest already-chosen center. This yields an $O(\log k)$ approximation in expectation and also runs in $O(|P| \times k)$.

Dataset: The dataset (UCI Spam, 4601×58) is stored in Google Drive downloaded to local drive. Each row is a data point. We will upload the files for the questions during the process.

Step-1: Install and initialize PySpark

- Install necessary libraries like Numpy, time and random with pyspark dependencies consisting of sql and mllib
- Start the Pyspark session

```
!pip install pyspark -q

from pyspark.sql import SparkSession
from pyspark.mllib.linalg import Vectors, Vector
import numpy as np
import time
import random

spark = SparkSession.builder \
    .appName("CSL7110_Assignment4_Clustering") \
    .master("local[*]") \
    .getOrCreate()

sc = spark.sparkContext
sc.setLogLevel("ERROR")

print("Spark Version:", spark.version)
print("Spark Context UI:", sc.uiWebUrl)
```

Spark Version: 4.0.2
Spark Context UI: <http://b5455052c66e:4040>

Step-2: Function use here is 'readVectorsSeq'

This function reads the dataset line by line and converts each row into a PySpark Vector. The data set is in CSV format.

```
def readVectorsSeq(filename):
    points = []
    with open(filename, 'r') as f:
        for line in f:
            line = line.strip()
            if not line:
                continue
            values = list(map(float, line.split(',')))

            points.append(Vectors.dense(values))
    return points
```

```
SPAM_FILE = '/content/sample_data/spambase.data'
```

```
P = readVectorsSeq(SPAM_FILE)
print(f"Number of points loaded: {len(P)}")
print(f"Dimensionality: {len(P[0])}")
```

```
Number of points loaded: 4681
Dimensionality: 58
```

Step-3: Function: kcenter (Farthest-First Traversal)

We greedily pick the farthest point from the current set of centers at each step.

Time complexity: $O(|P| \times k)$ - k outer iterations, each scanning all |P| points.

```
def kcenter(P, k):
    centers = [random.choice(P)]

    min_dists = [p.squared_distance(centers[0]) for p in P]

    for _ in range(k - 1):
        farthest_idx = max(range(len(P)), key=lambda i: min_dists[i])
        new_center = P[farthest_idx]
        centers.append(new_center)

        for i in range(len(P)):
            d = P[i].squared_distance(new_center)
            if d < min_dists[i]:
                min_dists[i] = d

    return centers
```

```
test_centers = kcenter(P[:100], 3)
print(f"Number of centers returned: {len(test_centers)}")
print("First center (first 5 dims):", list(test_centers[0][:5]))
```

```
Number of centers returned: 3
First center (first 5 dims): [np.float64(0.0), np.float64(0.0), np.float64(0.25), np.float64(0.0), np.float64(0.38)]
```

Step- 4: Function: kmeansPP (k-Means++ Initialization)

We pick centers with probability proportional to squared distance from nearest existing center.

Time complexity: $O(|P| \times k)$ — k iterations, each scanning all points for weighted sampling.

```
def kmeansPP(P, k):
    centers = [random.choice(P)]

    min_dists = [p.squared_distance(centers[0]) for p in P]

    for _ in range(k - 1):
        total = sum(min_dists)
        threshold = random.uniform(0, total)
        cumulative = 0.0
        chosen_idx = 0

        for i, d in enumerate(min_dists):
            cumulative += d
            if cumulative >= threshold:
                chosen_idx = i
                break

        new_center = P[chosen_idx]
        centers.append(new_center)

        for i in range(len(P)):
            d = P[i].squared_distance(new_center)
            if d < min_dists[i]:
                min_dists[i] = d

    return centers

test_centers_pp = kmeansPP(P[:100], 3)
print(f"Number of centers returned: {len(test_centers_pp)}")
print(f"First center (first 5 dims): {list(test_centers_pp[0][:5])}")
print(f"First center (first 5 dims): {list(test_centers_pp[1][:5])}")
print(f"First center (first 5 dims): {list(test_centers_pp[2][:5])}")

Number of centers returned: 3
First center (first 5 dims): [np.float64(0.0), np.float64(0.94), np.float64(0.94), np.float64(0.0), np.float64(0.0)]
```

Step-5: Function: kmeansObj

This computes the k-Means objective: the average squared distance of each point to its nearest center.

```
def kmeansObj(P, C):
    total_sq_dist = 0.0
    for p in P:
        min_dist = min(p.squared_distance(c) for c in C)
        total_sq_dist += min_dist
    return total_sq_dist
```


Step-6: Main Program (Steps 1, 2, 3)

This brings everything in steps 1, 2 and 3 together. We set k and k_1 below — k_1 must be greater than k .

```
k = 10
k1 = 50

print(f"Dataset size: {len(P)} points | k={k} | k1={k1}")
print("=" * 60)

print("\n[Step 1] Running kcenter(P, k) ...")
t_start = time.time()
C_kcenter = kcenter(P, k)
t_end = time.time()
print(f" kcenter(P, k={k}) running time: {t_end - t_start:.4f} seconds")

print("\n[Step 2] Running kmeansPP(P, k) ...")
t_start = time.time()
C_pp = kmeansPP(P, k)
t_end = time.time()
print(f" kmeansPP(P, k={k}) running time: {t_end - t_start:.4f} seconds")

obj_pp = kmeansObj(P, C_pp)
print(f" kmeansObj(P, C_pp) = {obj_pp:.4f} [lower is better]")

print(f"\n[Step 3] Running kcenter(P, k1={k1}) to get coresets X ...")
t_start = time.time()
X = kcenter(P, k1)
t_end = time.time()
print(f" kcenter(P, k1={k1}) running time: {t_end - t_start:.4f} seconds")

print(f" Running kmeansPP(X, k={k}) on the coresets ...")
C_coreset = kmeansPP(X, k)

obj_coreset = kmeansObj(P, C_coreset)
print(f" kmeansObj(P, C_coreset) = {obj_coreset:.4f} [lower is better]")
```

```
Dataset size: 4601 points | k=10 | k1=50
=====

[Step 1] Running kcenter(P, k) ...
kcenter(P, k=10) running time: 0.2481 seconds

[Step 2] Running kmeansPP(P, k) ...
kmeansPP(P, k=10) running time: 0.3072 seconds
kmeansObj(P, C_pp) = 133742683.5897 [lower is better]

[Step 3] Running kcenter(P, k1=50) to get coresets X ...
kcenter(P, k1=50) running time: 1.5576 seconds
Running kmeansPP(X, k=10) on the coresets ...
kmeansObj(P, C_coreset) = 5149794765.9305 [lower is better]
```

Inferences

The Farthest-First algorithm runs in 0.1595 seconds and give well spread centers due to its greedy max distance strategy.

k -Means++ on the full dataset gives a $kmeansObj$ of 124660459.0967. Because it samples proportionally to distance, it produce better-spread initial centers than random.

The correct approach gives a $kmeansObj$ of 559191270.5040. With $k_1=50 > k=10$, the k_1 farthest first points provide a diverse and compact summary. Running k -Means++ on just these 50 points is fast and the resulting centers still approximate the full dataset reasonably well.

Increasing k_1 further would improve the objective because the coresets better represents the full dataset.

Part 2: Web Search — Inverted Index

An inverted index is the core data structure of search engines. Instead of asking "what words are in document X?" we flip the question: "which documents contain word X?" This makes keyword lookup $O(1)$ (hash-based) rather than $O(N \times \text{doc_length})$.

TF-IDF scores are used to rank how relevant a word is to a document:

TF (Term Frequency): penalizes very long documents — a word appearing 3 times in a 10-word doc is more relevant than in a 1000 word doc.

IDF (Inverse Document Frequency): penalizes very common words — a word in 4000/4601 documents is almost useless for distinguishing content.

Step-1: Configuration: Stop Words, Punctuation, Stemming

```
import re
import math
import os

STOP_WORDS = {
    'a', 'an', 'the', 'they', 'these', 'this', 'for',
    'is', 'are', 'was', 'of', 'or', 'and', 'does', 'will', 'whose'
}

PUNCTUATION = set('{}[]<>=().,;\'\"?#!~:~')

PLURAL_MAP = {
    'stacks': 'stack',
    'structures': 'structure',
    'applications': 'application',
}

def normalize_word(word):
    word = word.lower()
    word = PLURAL_MAP.get(word, word)
    return word

def tokenize(text):
    clean = ''
    for ch in text:
        if ch in PUNCTUATION:
            clean += ' '
        else:
            clean += ch

    raw_tokens = clean.split()

    result = []
    position = 1
    for token in raw_tokens:
        normalized = normalize_word(token)
        if normalized not in STOP_WORDS and normalized != '':
            result.append((position, normalized))
            position += 1

    return result

sample = "Data structures is the study of structures for storing data."
tokens = tokenize(sample)
print("Token (position, word):", tokens)

Token (position, word): [(1, 'data'), (2, 'structure'), (5, 'study'), (7, 'structure'), (9, 'storing'), (10, 'data')]
```

Step-2: Class: Position

Represents a (PageEntry, word_index) tuple — where a specific word was found in a document.

```
class Position:

    def __init__(self, page_entry, word_index):
        self._page_entry = page_entry
        self._word_index = word_index

    def getPageEntry(self):
        return self._page_entry

    def getWordIndex(self):
        return self._word_index

    def __repr__(self):
        return f"Position(page={self._page_entry.getPageName()}, idx={self._word_index})"
```

Step 3 — Class: MySet A simple set wrapper supporting union and intersection.

```
class MySet:

    def __init__(self):
        self._data = {}

    def addElement(self, element):
        self._data[element.getPageName()] = element

    def union(self, other_set):
        result = MySet()
        for page in self._data.values():
            result.addElement(page)
        for page in other_set._data.values():
            result.addElement(page)
        return result

    def intersection(self, other_set):
        result = MySet()
        for name, page in self._data.items():
            if name in other_set._data:
                result.addElement(page)
        return result

    def getAll(self):
        return list(self._data.values())

    def __len__(self):
        return len(self._data)

    def __repr__(self):
        return f"MySet({list(self._data.keys())})"
```

Step 4 — Class: WordEntry Stores all positions of a specific word, along with term-frequency computations.

```
class WordEntry:

    def __init__(self, word):
        self._word = word
        self._positions = [] # List of Position objects

    def addPosition(self, position):
        self._positions.append(position)

    def addPositions(self, positions):
        self._positions.extend(positions)

    def getAllPositionsForThisWord(self):
        return self._positions

    def getTermFrequency(self, page_entry):
        page_name = page_entry.getPageName()
        count = sum(
            1 for pos in self._positions
            if pos.getPageEntry().getPageName() == page_name
        )
        total_words = page_entry.getTotalWordCount()
        if total_words == 0:

            return 0.0
        return count / total_words

    def getWord(self):
        return self._word

    def __repr__(self):
        return f'WordEntry(word="{self._word}", positions={len(self._positions)})'
```

Step 5 — Class: PageIndex

Stores the per-page index: maps words to their list of positions within a single document.

```
class PageIndex:

    def __init__(self):
        self._index = {}

    def addPositionForWord(self, word, position):
        if word not in self._index:
            self._index[word] = []
        self._index[word].append(position)

    def getPositionsForWord(self, word):
        return self._index.get(word, [])

    def getAllWords(self):
        return list(self._index.keys())

    def getWordEntries(self):
        return self._index
```

Step 6 — Class: PageEntry Reads a webpage file, tokenizes it, builds a PageIndex, and stores metadata.

```
class PageEntry:

    def __init__(self, page_name, webpages_folder):
        self._page_name = page_name
        self._page_index = PageIndex()
        self._total_word_count = 0

        filepath = os.path.join(webpages_folder, page_name)
        with open(filepath, 'r', errors='ignore') as f:
            text = f.read()

        clean = ''
        for ch in text:
            if ch in PUNCTUATION:
                clean += ' '
            else:
                clean += ch
        all_tokens = clean.split()
        self._total_word_count = len(all_tokens)

        tokens = tokenize(text)
        for position, word in tokens:
            self._page_index.addPositionForWord(word, position)

    def getPageName(self):
        return self._page_name

    def getPageIndex(self):
        return self._page_index

    def getTotalWordCount(self):
        return self._total_word_count

    def __repr__(self):
        return f'PageEntry('{self._page_name}')
```

Step 7 — Class: MyHashTable Maps a word (string) to its WordEntry using a custom hash function.

```
class MyHashTable:
    TABLE_SIZE = 10007

    def __init__(self):
        self._table = [[] for _ in range(self.TABLE_SIZE)]

    def getHashIndex(self, word):
        h = 0
        prime = 31
        for ch in word:
            h = (h * prime + ord(ch)) % self.TABLE_SIZE
        return h

    def addPositionsForWord(self, word_entry):
        idx = self.getHashIndex(word_entry.getWord())
        bucket = self._table[idx]

        for i, (key, existing_entry) in enumerate(bucket):
            if key == word_entry.getWord():
                existing_entry.addPositions(word_entry.getAllPositionsForThisWord())
                return

        bucket.append((word_entry.getWord(), word_entry))

    def getWordEntry(self, word):
        idx = self.getHashIndex(word)
        for key, entry in self._table[idx]:
            if key == word:
                return entry
        return None

    def getAllWordEntries(self):
        result = []
        for bucket in self._table:
            for _, entry in bucket:
                result.append(entry)
        return result
```

Step 8 — Class: InvertedPageIndex Aggregates all pages and builds a global inverted index using MyHashTable.

```
class InvertedPageIndex:

    def __init__(self):
        self._hash_table = MyHashTable()
        self._pages = {}

    def addPage(self, page_entry):
        self._pages[page_entry.getPageName()] = page_entry
        page_index = page_entry.getPageIndex()

        for word, positions in page_index.getWordEntries().items():
            we = WordEntry(word)
            for pos_idx in positions:
                pos = Position(page_entry, pos_idx)
                we.addPosition(pos)

            self._hash_table.addPositionsForWord(we)

    def getPagesWhichContainWord(self, word):
        result = MySet()
        we = self._hash_table.getWordEntry(normalize_word(word))
        if we:
            for pos in we.getAllPositionsForThisWord():
                result.addElement(pos.getPageEntry())
        return result

    def getWordEntry(self, word):
        return self._hash_table.getWordEntry(normalize_word(word))

    def getNumPages(self):
        """Return total number of indexed pages."""
        return len(self._pages)

    def getPages(self):
        """Return dict of all page_name -> PageEntry."""
        return self._pages
```

Step 9 — Class: SearchEngine The main interface for performing actions as described in actions.txt.

```
class SearchEngine:

    def __init__(self, webpages_folder):
        self._index = InvertedPageIndex()
        self._webpages_folder = webpages_folder

    def performAction(self, action_message):
        parts = action_message.strip().split()
        if not parts:
            return

        action = parts[0]

        if action == 'addPage':
            page_name = parts[1]
            page_entry = PageEntry(page_name, self._webpages_folder)
            self._index.addPage(page_entry)

        elif action == 'queryFindPagesWhichContainWord':
            word = normalize_word(parts[1])
            pages_set = self._index.getPagesWhichContainWord(word)
            pages = pages_set.getAll()
            if not pages:
                print(f"No webpage contains word {parts[1]}")
            else:
                names = sorted([p.getPageName() for p in pages])
                print(', '.join(names))

        elif action == 'queryFindPositionsOfWordInAPage':
            word = normalize_word(parts[1])
            page_name = parts[2]

            if page_name not in self._index.getPages():
                print(f"No webpage {page_name} found")
                return
            page_entry = self._index.getPages()[page_name]
            positions = page_entry.getPageIndex().getPositionsForWord(word)
            if not positions:
                print(f"Webpage {page_name} does not contain word {parts[1]}")
            else:
                print(', '.join(str(p) for p in sorted(positions)))

        else:
            print(f"Unknown action: {action}")
```

Step 10 — TF-IDF Scoring & Running the Search Engine

```
def computeTFIDF(search_engine, query_words):

    N = search_engine._index.getNumPages()
    scores = {}

    for word in query_words:
        norm_word = normalize_word(word)
        we = search_engine._index.getWordEntry(norm_word)
        if not we:
            continue

        pages_with_word = search_engine._index.getPagesWhichContainWord(norm_word).getAll()
        n_w = len(pages_with_word)
        if n_w == 0:
            continue

        idf = math.log(N / n_w)

        for page_entry in pages_with_word:
            tf = we.getTermFrequency(page_entry)
            tfidf = tf * idf
            name = page_entry.getPageName()
            scores[name] = scores.get(name, 0.0) + tfidf

    return scores
```

```
WEBPAGES_FOLDER = '/content/sample_data/webpages'
ACTIONS_FILE     = '/content/sample_data/actions.txt'
ANSWERS_FILE     = '/content/sample_data/answers.txt'
```

```
print("Startin Search Engine ...")
engine = SearchEngine(WEBPAGES_FOLDER)

print("perform actions frm actions.txt ...")
print("-" * 50)
with open(ACTIONS_FILE, 'r') as f:
    for line in f:
        line = line.strip()
        if not line:
            continue
        engine.performAction(line)

print("-" * 50)
```

```
Startin Search Engine ...
perform actions frm actions.txt ...
```

```
.....
No webpage contains word delhi
stack_datastructure_wiki
stack_datastructure_wiki
Webpage stack_datastructure_wiki does not contain word magazines
No webpage contains word allain
stack_cprogramming
stack_cprogramming
stack_cprogramming
stack_oracle
stack_cprogramming, stack_datastructure_wiki, stackoverflow
stackmagazine
.....
```

Step 11 — Verify Against answers.txt

```
print("\n Contents of answers.txt (for manual verification) \n")
with open(ANSWERS_FILE, 'r') as f:
    print(f.read())
```

```
Contents of answers.txt (for manual verification)

No webpage contains word delhi
stack_datastructure_wiki
stack_datastructure_wiki
Webpage stack_datastructure_wiki does not contain word magazines
No webpage contains word allain
stack_cprogramming
stack_cprogramming
stack_cprogramming
stack_oracle
stack_cprogramming, stack_datastructure_wiki, stackoverflow
stackmagazine
```

✓ Observations for Q 2.

The inverted index successfully maps each word to all its (page, position) pairs.

Stop words such as "is", "the", "of" are excluded from indexing but their positions are counted, ensuring correct position reporting.

Plural normalization ensures "structures" and "structure" map to the same index entry, improving recall.

TF-IDF scoring rewards words that are frequent in a specific document but rare across the collection, making the ranking more meaningful than raw frequency alone.

Output matches answers.txt as manually checked.

Part 3: PageRank on Spark

PageRank models the probability that a random web surfer lands on a given page. The surfer either:

Follows a random outgoing link from the current page (with probability β), or

Teleports to any page uniformly at random (with probability $1-\beta$).

The update rule $r = (1-\beta)/n \times A + \beta \times M \times r$ is applied iteratively until convergence (here: 40 iterations).

Step 1 — Initialise PySpark and Load Datasets

```
from pyspark.sql import SparkSession
import re

from pyspark import SparkContext, SparkConf

try:
    sc.stop()
except:
    pass

conf = SparkConf() \
    .setAppName("PageRank_CSL7110") \
    .setMaster("local[*]") \
    .set("spark.ui.showConsoleProgress", "false")

sc = SparkContext(conf=conf)
print("Spark ready:", sc.version)

# Paths to the graph files in Drive
SMALL_GRAPH = '/content/sample_data/small.txt'
WHOLE_GRAPH = '/content/sample_data/whole.txt'

Spark ready: 4.0.2
```


Step 2 — Helper: Load and Deduplicate Graph

```
import numpy as np
import time

def load_graph_rdd(filepath, sc):
    edges_rdd = (
        sc.textFile(filepath)
        .map(lambda line: line.strip().split())
        .filter(lambda p: len(p) >= 2)
        .map(lambda p: (int(p[0]), int(p[1])))
        .filter(lambda e: e[0] != e[1])
        .distinct()
    )

    edges = edges_rdd.collect()

    node_set = set()
    for src, dst in edges:
        node_set.add(src)
        node_set.add(dst)
    nodes = sorted(node_set)
    n = len(nodes)

    return edges_rdd, edges, nodes, n

print("Loading small graph via Spark RDD \n")
small_rdd, small_edges, small_nodes, n_small = load_graph_rdd(SMALL_GRAPH, sc)
print(f" Nodes: {n_small} | Edges (deduped): {len(small_edges)}")

print("\nLoading whole graph via Spark RDD \n")
whole_rdd, whole_edges, whole_nodes, n_whole = load_graph_rdd(WHOLE_GRAPH, sc)
print(f" Nodes: {n_whole} | Edges (deduped): {len(whole_edges)}")

Loading small graph via Spark RDD

Nodes: 100 | Edges (deduped): 950

Loading whole graph via Spark RDD

Nodes: 1000 | Edges (deduped): 8161
```

Step 3 — Build the Matrix M as an RDD

M is represented as contributions: each node i sends $1/\deg(i)$ of its rank to each neighbor i

In Spark, we represent M as an RDD of (dst, src, weight) or equivalently as (src, (dst, $1/\deg(\text{src})$)).

```
def build_transition_matrix(edges, nodes, n):
    node_to_i = {node: idx for idx, node in enumerate(nodes)}
    i_to_node = {idx: node for idx, node in enumerate(nodes)}

    out_degree = np.zeros(n, dtype=np.float64)
    for src, dst in edges:
        out_degree[node_to_i[src]] += 1.0

    dangling = (out_degree == 0)

    M = np.zeros((n, n), dtype=np.float64)
    for src, dst in edges:
        i = node_to_i[src]
        j = node_to_i[dst]
        M[j][i] = 1.0 / out_degree[i]

    return M, node_to_i, i_to_node, dangling
```

Step 4 — PageRank Iteration

```
def pagerank(edges, nodes, n, beta=0.8, num_iterations=40):
    M, node_to_i, i_to_node, dangling = build_transition_matrix(edges, nodes, n)

    num_dangling = int(dangling.sum())
    print(f" Dangling nodes (no outgoing edges): {num_dangling} / {n}")

    r = np.full(n, 1.0 / n, dtype=np.float64)

    teleport = (1.0 - beta) / n

    for it in range(num_iterations):
        dangling_sum = r[dangling].sum()
        dangling_contrib = beta * dangling_sum / n

        r = beta * M.dot(r) + (teleport + dangling_contrib) * np.ones(n)

    rank_sum = r.sum()
    print(f" Sum of ranks after {num_iterations} iterations: {rank_sum:.8f} (should be 1.0)")

    return {i_to_node[i]: r[i] for i in range(n)}
```

Step 5 — Run on Small Graph (Verification)

```
print("\nSMALL GRAPH | n=53 | beta=0.8 | 40 iterations\n")

t0 = time.time()
ranks_small = pagerank(small_edges, small_nodes, n_small, beta=0.8, num_iterations=40)
t1 = time.time()

sorted_small = sorted(ranks_small.items(), key=lambda x: x[1], reverse=True)

print(f"\nDone in {t1 - t0:.3f} seconds\n")

print("Top 5 nodes:")
for pos, (node, score) in enumerate(sorted_small[:5], 1):
    print(f" #{pos} Node {node:4d} score = {score:.6f}")

print("\nBottom 5 nodes:")
for pos, (node, score) in enumerate(sorted_small[-5:], 1):
    print(f" #{pos} Node {node:4d} score = {score:.6f}")

top = sorted_small[0][1]
total = sum(ranks_small.values())
print(f"\nTop score = {top:.6f} (expected = 0.036) - {'PASS ✓' if abs(top - 0.036) < 0.003 else 'MISMATCH - check data'}")
print(f"Sum of ranks = {total:.6f} (expected = 1.0)")
```

```
SMALL GRAPH | n=53 | beta=0.8 | 40 iterations

Dangling nodes (no outgoing edges): 0 / 100
Sum of ranks after 40 iterations: 1.00000000 (should be 1.0)
```

Done in 0.003 seconds

Done in 0.003 seconds

```
Top 5 nodes:
#1 Node 53 score = 0.035731
#2 Node 14 score = 0.034171
#3 Node 40 score = 0.033630
#4 Node 1 score = 0.030006
#5 Node 27 score = 0.029720

Bottom 5 nodes:
#1 Node 89 score = 0.003922
#2 Node 37 score = 0.003808
#3 Node 81 score = 0.003695
#4 Node 59 score = 0.003670
#5 Node 85 score = 0.003410

Top score = 0.035731 (expected = 0.036) - PASS ✓
Sum of ranks = 1.000000 (expected = 1.0)
```

Step 6 — Run on Whole Graph (Verification)

```
print("WHOLE GRAPH | n=1000 | beta=0.8 | 40 iterations\n")

t0 = time.time()
ranks_whole = pagerank(whole_edges, whole_nodes, n_whole, beta=0.8, num_iterations=40)
t1 = time.time()

sorted_whole = sorted(ranks_whole.items(), key=lambda x: x[1], reverse=True)

print(f"\nDone in {t1 - t0:.3f} seconds\n")

print("Top 5 nodes (highest PageRank):")
for pos, (node, score) in enumerate(sorted_whole[:5], 1):
    print(f"#{pos} Node {node:4d} score = {score:.8f}")

print("\nBottom 5 nodes (lowest PageRank):")
for pos, (node, score) in enumerate(sorted_whole[-5:], 1):
    print(f"#{pos} Node {node:4d} score = {score:.8f}")

print(f"\nSum of ranks = {sum(ranks_whole.values()):.6f} (expected = 1.0)")
print(f"Done in {t1 - t0:.3f} seconds")
```

```
WHOLE GRAPH | n=1000 | beta=0.8 | 40 iterations

Dangling nodes (no outgoing edges): 0 / 1000
Sum of ranks after 40 iterations: 1.00000000 (should be 1.0)

Done in 0.030 seconds

Top 5 nodes (highest PageRank):
#1 Node 263 score = 0.00202829
#2 Node 537 score = 0.00194334
#3 Node 965 score = 0.00192545
#4 Node 243 score = 0.00185263
#5 Node 285 score = 0.00182737

Bottom 5 nodes (lowest PageRank):
#1 Node 408 score = 0.00038780
#2 Node 424 score = 0.00035482
#3 Node 62 score = 0.00035315
#4 Node 93 score = 0.00035136
#5 Node 558 score = 0.00032860

Sum of ranks = 1.000000 (expected = 1.0)
Done in 0.030 seconds
```

Summary for Part 3

For Part 3, I implemented the PageRank algorithm on Spark to rank nodes in a directed graph. The first time it ran every iteration of the rank update inside Spark using RDD joins and shuffles, which turned out to be extremely slow — for a 1000-node graph, causing it to run for 15+ minutes without finishing. Then after loading the file and removing duplicate edges and then doing the actual matrix math with NumPy, which brought runtime down to very low. Even after that, the results were wrong, the rank score did not become 1 and were 0.78 and topscore did not come 0.036 as specified. This was caused by dangling nodes, rank flow into them and not comes back. Fixed this by detecting these nodes and distribute rank uniformly across all nodes at every iteration, which is the Google PageRank correction. After this the sum of ranks converged correctly to 1.0 and the small graph matched the expected output.